

Trends and developments in Reinforcement Learning

London TensorFlow Meetup

Pierre H. Richemond

Imperial College London

Table of contents

1. Introduction: the RL Setting
2. Theoretical Developments
3. Progress in continuous control
4. Progress in board games
5. Some practical tips

Introduction: the RL Setting

RL in one minute !

- Reinforcement learning = adaptive, extreme trial-and-error
- Deep RL = RL combined with deep neural networks
- Given a certain **environment** (game abstraction), an **agent** (player) selects between **actions** a available to her, and in doing so, collects a **reward** r and moves from **state** s to state s' .
- The agent's objective is to pick actions and build a trajectory so as to **maximize the sum of their rewards**.
- Rewards and states are not known in advance, so that RL involves a trade-off between **exploration** of new and risky states and **exploitation** of known but potentially sub-optimal rewards.

Just one more minute :)

Two main families of algorithmic methods for solving RL:

1. **Q-learning** consists in computing an approximation to an *action-value function* $Q(s, a)$ which for each state-action pair tells us how happy we will be if we take that action in that state.
2. **Policy gradients** consists in approximately following the gradient to the sum of rewards ('do more of what works best') so as to maximize it. It seeks to compute $\pi(a|s) = \mathbb{P}(a|s)$.

Deep Q-Networks (DQN) and Asynchronous Advantage Actor Critic (A3C) are the most famous deep RL algorithms.

In general, RL methods are **numerically unstable** and require tricks to stabilize them and accelerate their convergence.

Theoretical Developments

Equivalence of soft Q-learning and policy gradients

- The single most important theoretical development of 2017.
- We've been doing reinforcement learning wrong all these years !
- Idea = turn the **hard** max in Q-learning into a **soft** max.
- Mathematically, relax $Q^*(s, a) = \max_{a \in \mathbb{A}} Q(s, a)$ into

$$Q^*(s, a) = \beta \log \int \exp\left(\frac{Q(s, a)}{\beta}\right) da$$

- This relaxation was already known since 2010 at least.

Equivalence of soft Q-learning and policy gradients (2)

- Policy gradients often use **entropy regularization** to prevent early convergence to degenerate policies.
- Schulman shows that in this context, soft Q-learning and entropic policy gradients are **the same**.
- Easily proven with an entropic representation formula (Richemond & Maginnis, to appear.)
- This also proves that the A3C algorithm as originally published is wrong/incomplete and explains why it may fail to converge.
- Strong theoretical connections both with thermodynamics and convex/online optimization.

Wasserstein RL solves the heat equation

- New development, on ArXiv soon
- Optimal transport / *Wasserstein* distances have become popular in the context of GANs
- When doing trust region policy optimization in a small Wasserstein-2 ball rather than Kullback-Leibler, one proves that the policy function itself, in the small step limit, satisfies the heat equation

$$\partial_t \pi = -\nabla \cdot (\pi \nabla r) + \beta \Delta \pi$$

- This might also be a justification for the emergence of *distributional* RL

Progress in continuous control

Proximal Policy Optimization - Robust knocked over stand up

<https://www.youtube.com/watch?v=bqdjsmSoSgI>

Emergence of Locomotion Behaviours in Rich Environments

https://www.youtube.com/watch?v=hx_bg0TF7bs&t=49s

- *Trust Region Policy Optimization* is a very robust policy gradient baseline that iteratively refines policies by moving them within a small stepsize.
- This approach introduces strong sensitivity w.r.t. the stepsize as an important tunable hyperparameter and makes it generally sample inefficient
- *Proximal Policy Optimization* is an evolution of TRPO that attempts to do away with those constraints
- The idea is to use a new objective function that includes clipping the ratio of the probability under the new and old policies
- Simple change in the code.

- Combines actor-critic, trust region optimization, and Kronecker factorization ('K-FAC')
- As this is a second-order (*natural gradient*) method, it's more sample efficient
- Flipside is increased computational complexity
- Scales well with batch size
- Together PPO and ACKTR represent the state-of-the-art baseline in continuous control... until the official release of *soft actor-critic* (2018, OpenAI)

Progress in board games

- Arguably the landmark practical achievement of the year in reinforcement learning !
- AlphaGo the first learned the game by analyzing 30 million Go moves from human players (*GoKiFu* dataset) before refining that supervised policy with self-play
- This illustrates the difficulty of the mathematical optimization problem posed by RL, and the necessity to initialize close to a good solution.
- AlphaGo Zero plays significantly better, and does away with supervision or human knowledge completely (other than knowledge of game rules)

AlphaGo Zero: Architecture (1)

- Problem with Go = large branching factor in the tree of possible moves.
- Neural networks enable dynamic pruning of that tree - a *policy* network limits the breadth of the subtree to explore, while a *value* network limits the depth of that subtree. *If those networks were perfect, little exploration would be needed - in order to get there, we can refine them iteratively.*
- Merging the policy and value network enforces feature reuse and learning efficiency
- In order to train deeper networks and enable more and more abstract feature extraction, *residual networks* are used

AlphaGo Zero: Architecture (2)

- **Monte Carlo Tree Search** (MCTS) is another part of the algorithm that, just like a neural network, dynamically expands and compresses the tree of possibles. It averages results across many paths of the sub-tree, and can be viewed as a form of imagination, or slow thinking.
- The key component of AlphaGo Zero is to view MCTS as a *policy improvement* mechanism: you can't improve anymore when your fast thinking (NN) and your slow thinking (MCTS) coincide.
- The loss function is "simple" for a program that complex:

$$L(\theta) = (z - v_\theta)^2 + \pi^T \log p_\theta + c \|\theta\|^2$$

AlphaGo Zero: Insights

- *AlphaZero* is the generic, game-agnostic version of AlphaGoZero that also taught itself to play Chess and Shogi (Japanese chess where you can capture your opponent's pieces) at a superhuman level using the same algorithm.
- In stark contrast with DeepBlue or Stockfish, AlphaZero 'only' evaluates 50,000 positions per move due to MCTS, rather than millions.
- AlphaGo has now become so strong it is used as an online teaching tool for human players - *AlphaGo Teach*
<https://alphagoteach.deepmind.com/>
- According to David Silver, it is speculated that MCTS acts as a large averaging mechanism, rather than the 'maximum' operation in standard tree pruning, **smoothing rather than propagating** neural network approximation error, and hence **enabling stable training** of deep networks.

Some practical tips

Suggestions for experiment design

Factors influencing training, by order of importance:

- **Fix your random seed** for reproducibility !
- *Softmax temperature* is the single most critical factor according to Nachum & Huarnooja, followed by...
- *Reward scaling* and more generally reward engineering help a lot
- *Number of steps* in the calculation of returns
- *Discount factor* also an issue

Clip gradients to avoid NaNs, visualize them with TensorBoard, and finally be **patient** waiting for convergence...

Key TensorFlow code repos

- As close to official as possible :
https://github.com/tensorflow/models/tree/master/research/pcl_rl
by Ofir Nachum the author of PCL, soon to be integrated in TensorFlow
- RLlib: A Composable and Scalable Reinforcement Learning Library
<https://github.com/ray-project/ray/tree/master/python/ray/rllib>
- Imperial's Brendan Maginnis Atari-RL
<https://github.com/brendanator/atari-rl>
- OpenAI's *baselines*
<https://github.com/openai/baselines>

Deep Reinforcement Learning That Matters

- Encouraging reproducibility in reinforcement learning
- 'RL algorithms have many moving parts that are hard to debug, and they require substantial effort in tuning in order to get good results.'
- False Discoveries everywhere ? Dependency on the random seed
- Multiple runs are necessary to give an idea of the variability. Reporting the best returns is not enough - at the very least, report top and bottom quantiles
- A necessary debate, as seen at NIPS 2017.

Don't be an alchemist

Conclusion & perspectives

- Another landmark year both for theoretical and practical RL
- In imperfect information settings - Poker : *Libratus* and *DeepStack* become superhuman at Texas Hold'em
- In robotics, notable progress also in *hierarchical reinforcement learning*, *meta-learning* and *imitation learning*.
- Planning, robustness and sample efficiency remain the greatest challenges to current methods
- The next big challenge : Starcraft II, a real-time strategy game, partially observed, with sparse rewards and a huge action space. PySC2 available to train agents at

<https://github.com/deepmind/pysc2>

Questions from the audience ?

Thanks for your time !

Twitter @KloudStrife